

Creating EC2 Instances

Within the scope of the project, a virtual machine will be created to install nginx web server. As virtual machine features:

1. Amazon AMI 2 will be used as it is free and resource sufficient.
2. t3.micro instance type will be used. It will provide a sufficient and free configuration for the server.
3. Vpc and security groups created previously in the Network tab will be used. Also Auto-Assign public IP option is selected as enable. This option will automatically assign an IP address to the EC2 server.
4. In the Advanced Options tab, the following code line will be written in "User Data".

```
#!/bin/bash
```

```
yum update -y # Update operating system packages
```

```
# Install Nginx web server
```

```
yum install nginx -y
```

```
yum upgrade -y # Update all operating system packages
```

```
yum autoremove -y # Remove unused packages
```

```
yum clean all # Clear cache
```

```
systemctl enable nginx # Start the Nginx service automatically
```

```
systemctl start nginx # Start the Nginx service
```

Launch template name and description

Launch template name - *required*

Nginx-Web-Server

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description

Nginx Web Server.

Max 255 chars

Auto Scaling guidance [Info](#)

Select this if you intend to use this template with EC2 Auto Scaling

☐ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

► Template tags


► Source template

▼ Application and OS Images (Amazon Machine Image) - required [Info](#)


An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

🔍 Search our full catalog including 1000s of application and OS images


Quick Start




Amazon Linux




macOS




Ubuntu




Windows



Red Hat



SUSE Linux



[Browse more AMIs](#)

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type
ami-0617f77bd30aaddb7 (64-bit (x86)) / ami-062725e4470c36327 (64-bit (Arm))
Virtualization: hvm ENA enabled: true Root device type: ebs

Free tier eligible ▼

Description

Amazon Linux 2 Kernel 5.10 AMI 2.0.20230628.0 x86_64 HVM gp2

Architecture

64-bit (x86) ▼

AMI ID

ami-0617f77bd30aaddb7

Verified provider

▼ Instance type [Info](#)

[Advanced](#)

Instance type

t3.micro

Free tier eligible

Family: t3 2 vCPU 1 GiB Memory Current generation: true
On-Demand RHEL pricing: 0.0708 USD per Hour
On-Demand SUSE pricing: 0.0108 USD per Hour
On-Demand Linux pricing: 0.0108 USD per Hour
On-Demand Windows pricing: 0.02 USD per Hour

☒ All generations

[Compare instance types](#)

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name

Nginx Web Server

[Create new key pair](#)

▼ Network settings [Info](#)

Subnet [Info](#)

subnet-057e144bfc45d349d

Nginx Web Subnet

VPC: vpc-097825383c424a849 Owner: 273737642040
Availability Zone: eu-north-1a IP addresses available: 251 CIDR: 10.0.0.0/24

[Create new subnet](#)

When you specify a subnet, a network interface is automatically added to your template.

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Select existing security group

☐ Create security group

Common security groups [Info](#)

Select security groups

Nginx Security Group sg-009c73d747cfb98fb ✕
VPC: vpc-097825383c424a849

[Compare security group rules](#)

Security groups that you add or remove here will be added to or removed from all your network interfaces.

► [Advanced network configuration](#)

User data - *optional* [Info](#)

Upload a file with your user data or enter it in the field.

 Choose file

```
#!/bin/bash
yum update -y
yum install nginx -y
yum upgrade -y
yum autoremove -y
yum clean all
systemctl enable nginx
systemctl start nginx
```

☐ User data has already been base64 encoded

Creating Auto Scaling Group

Auto scaling means that aws services can be scaled automatically. This means that the memory of EC2 servers can be increased or decreased as needed. Auto scaling can also automatically deploy EC2 instances running in multiple Availability Zones to ensure high availability. This ensures that your application runs without interruption and is always available to users.

In auto scaling, the ec2 launch template that we previously created as launch template will be used. In the Network tab, vpc and subnets created for the project will be used. While creating the network in the project, only 1 subnet was created. To make an application running in more than one availability zone, one more subnet is created here.

Launch template [Info](#)

i For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template

Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

Nginx-Web-Server 

[Create a launch template](#)

Version

Default (1) 

[Create a launch template version](#)

Description

Nginx Web Server.

Launch template

[Nginx-Web-Server](#)

lt-03cf500d63db84e59

Instance type

t3.micro

AMI ID

ami-0617f77bd30aaddb7

Security groups

-

Request Spot Instances

No

Key pair name

Nginx Web Server

Security group IDs

[sg-009c73d747cfb98fb](#)

Additional details

Storage (volumes)

-

Date created

Wed Jul 19 2023 15:08:40
GMT+0300 (GMT+03:00)

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.


vpc-097825383c424a849 (Nginx Web Server) 
10.0.0.0/16


[Create a VPC](#)

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets 

eu-north-1a | subnet-0f285dce115e66052 (Nginx Web Server) 
10.0.0.0/24

eu-north-1b | subnet-0e0c87f9300987cab (Nginx Server Different AZ) 
10.0.1.0/24

[Create a subnet](#)

Instance type requirements [Info](#)

[Override launch template](#)

You can keep the same instance attributes or instance type from your launch template, or you can choose to override the launch template by specifying different instance attributes or manually adding instance types.

Launch template

[Nginx-Web-Server](#)

lt-03cf500d63db84e59

Version

Default

Description

Nginx Web Server.

Instance type

t3.micro

Cancel

Skip to review

Previous

Next

The load balancer receives incoming requests and distributes these requests to multiple servers or resources according to a specific algorithm or method. In this way, each server and resource are loaded evenly, performance is increased, and load balancing is performed.

In the Load balancer tab, the "Attach to a new Load Balancer" option was selected because we want to add a new load balancer. In the Attach to a new load balancer tab, the "Application Load Balancer" (ALB) option was selected. ALB is a layer 7 (application layer) load balancer used to route HTTP and HTTPS traffic. It has features such as application-level routing, dynamic scaling, multi-protocol support and advanced monitoring. "Internet-facing": This option indicates that your load balancer will be open to the internet. That is, it accepts requests from outside and directs these requests to the correct destinations. It is a preferred option for services open to the internet such as web applications.

Load balancing [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☐ Attach to an existing load balancer
Choose from your existing load balancers.

☒ Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to a new load balancer
Define a new load balancer to create for attachment to this Auto Scaling group.

Load balancer type
Choose from the load balancer types offered below. Type selection cannot be changed after the load balancer is created. If you need a different type of load balancer than those offered here, visit the [Load Balancing console](#).

☒ Application Load Balancer
HTTP, HTTPS

☐ Network Load Balancer
TCP, UDP, TLS

Load balancer name
Name cannot be changed after the load balancer is created.

Ngix Web Server-1

Load balancer scheme
Scheme cannot be changed after the load balancer is created.

☐ Internal

☒ Internet-facing

Network mapping
Your new load balancer will be created using the same VPC and Availability Zone selections as your Auto Scaling group. You can select different subnets and add subnets from additional Availability Zones.

VPC
vpc-097825383c424a849 [↗](#) Ngix Web Server

Availability Zones and subnets
You must select a single subnet for each Availability Zone enabled. Only public subnets are available for selection to support DNS resolution.

☒ eu-north-1b

subnet-0e0c87f9300987cab ▼

☒ eu-north-1a

subnet-0f285dce115e66052 ▼

Scaling Policies are directives that enable AWS Auto Scaling to automatically perform scaling operations based on a specific rule. These policies allow you to set your scaling goals and ensure that your application scales automatically based on demand.

By selecting the "Target tracking scaling policy" option in the scaling policies tab, it will automatically add and remove capacity to reach a specific target. Since the rule will be created according to CPU usage in the scaling policy name option, the name "CPUUtilisationScalingPolicy" is used as the name (Different name can be used if desired). In the metric type option, it is specified that the rule will be created according to the CPU usage of the EC2 instance by selecting the "Average CPU Utilisation" option. Selecting Target Value: 50% and instances need: 300; It is stated that a new instance will be created when the EC2 instance exceeds 50% utilisation for 300 seconds.

The screenshot shows the 'Scaling policies - optional' configuration page in the AWS IAM console. It includes a section to choose a scaling policy, with 'Target tracking scaling policy' selected. Below this, the 'Scaling policy name' is set to 'CPUUtilizationScalingPolicy', the 'Metric type' is 'Average CPU utilization', the 'Target value' is '50', and 'Instances need' is '300' seconds. There is also a checkbox for 'Disable scale in to create only a scale-out policy' which is currently unchecked.

Scaling policies - optional

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand. [Info](#)

☒ **Target tracking scaling policy**
Choose a desired outcome and leave it to the scaling policy to add and remove capacity as needed to achieve that outcome.

☐ None

Scaling policy name
CPUUtilizationScalingPolicy

Metric type
Average CPU utilization ▼

Target value
50

Instances need
300 seconds warm up before including in metric

☐ Disable scale in to create only a scale-out policy

A new sns has been created in the Add notification tab. AWS SNS (Simple Notification Service) is a service used to provide instant communication and notifications between your applications and systems. SNS allows you to send push notifications to your subscribers (message queues, email, SMS, HTTP endpoints, AWS Lambda functions, etc.) when an event occurs, or a certain condition is met.

The screenshot shows the 'Add notification' configuration page in the AWS IAM console. It includes a section for 'Notification 1' with a 'Remove' button. The 'Send a notification to' field is set to 'Nginx-Project-SNS-Topic'. The 'With these recipients' field is set to 'xxx@gmail.com'. There is a button for 'Use existing topic'. The 'Event types' section lists 'Launch', 'Terminate', 'Fail to launch', and 'Fail to terminate', all of which are checked. At the bottom, there is an 'Add notification' button.

▼ **Notification 1** Remove

Send a notification to
Nginx-Project-SNS-Topic

With these recipients
xxx@gmail.com

Use existing topic

Event types
Notify subscribers whenever instances

- ☒ Launch
- ☒ Terminate
- ☒ Fail to launch
- ☒ Fail to terminate

Add notification

TROUBLESHOOTING

1. SSH connection was not available. I thought that the reason for this was that the public IP address was not assigned, so I first tried to get an Elastic IP address and associate this address for the required instance. I thought it would be fixed when I rebooted the instance, but it was not. For this I ran network track analysis (also known as run reachability analysis). Here it was realized that the problem was in the route table. It seemed that no route was created for the internet gateway in the route table. For this, a new route was created in the route table:
Destination: 0.0.0.0/0 and target: igw-igw_ID. In this configuration, all traffic with the destination 0.0.0.0/0 is routed to the Internet Gateway specified as "igw-igw_ID". In this way, private subnets in the VPC can access the internet and, for example, EC2 instances can go to the internet or respond to requests from the internet.
2. The commands written in the User Data field were not executed because the server is not rooted. For this, the commands were executed again after connecting with ssh. The nginx configuration command was rewritten as: ***sudo amazon-linux-extras install nginx1.12 -y***.
3. While creating auto scaling, an error was encountered on the load balancer side, so the load balancer was not created. For this, a new load balancer was created from the load balancer tab and the load balancer was configured by connecting to the required vpc and subnets.

VPC > Route tables > rtb-0da8a3182c3a419b7 > Edit routes

Edit routes

Destination	Target	Status	Propagated
10.0.0.0/16	local	Active	No
0.0.0.0/0	igw-0bf0efdc0a1a375ce	Active	No

[Add route](#) [Cancel](#) [Preview](#) [Save changes](#)

Basic configuration

Load balancer name
Name must be unique within your AWS account and can't be changed after the load balancer is created.

Nginx-Load-Balancer

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme [Info](#)
Scheme can't be changed after the load balancer is created.

☒ Internet-facing
An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. [Learn more](#)

☐ Internal
An internal load balancer routes requests from clients to targets using private IP addresses.

IP address type [Info](#)
Select the type of IP addresses that your subnets use.

☒ IPv4
Recommended for internal load balancers.

☐ Dualstack
Includes IPv4 and IPv6 addresses.

Network mapping [Info](#)

The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.

VPC [Info](#)

Select the virtual private cloud (VPC) for your targets or you can [create a new VPC](#). Only VPCs with an internet gateway are enabled for selection. The selected VPC can't be changed after the load balancer is created. To confirm the VPC for your targets, view your [target groups](#).

Nginx Web Server

vpc-097825383c424a849

IPv4: 10.0.0.0/16



Mappings [Info](#)

Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.

☒ eu-north-1a (eun1-az1)

Subnet

subnet-0f285dce115e66052

Nginx Web Server ▼

IPv4 address

Assigned by AWS

☒ eu-north-1b (eun1-az2)

Subnet

subnet-0e0c87f9300987cab

Nginx Server Different AZ ▼

IPv4 address

Assigned by AWS

Basic configuration

Settings in this section can't be changed after the target group is created.

Choose a target type

☒ Instances

- Supports load balancing to instances within a specific VPC.
- Facilitates the use of [Amazon EC2 Auto Scaling](#) to manage and scale your EC2 capacity.

☐ IP addresses

- Supports load balancing to VPC and on-premises resources.
- Facilitates routing to multiple IP addresses and network interfaces on the same instance.
- Offers flexibility with microservice based architectures, simplifying inter-application communication.
- Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.

☐ Lambda function

- Facilitates routing to a single Lambda function.
- Accessible to Application Load Balancers only.

☐ Application Load Balancer

- Offers the flexibility for a Network Load Balancer to accept and route TCP requests within a specific VPC.
- Facilitates using static IP addresses and PrivateLink with an Application Load Balancer.

Target group name

Nginx-Target-Group

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Protocol

HTTP ▼

Port

: 80

1-65535

VPC

Select the VPC with the instances that you want to include in the target group.

Nginx Web Server

vpc-097825383c424a849

IPv4: 10.0.0.0/16

Protocol version

☐ HTTP1

Send requests to targets using HTTP/1.1. Supported when the request protocol is HTTP/1.1 or HTTP/2.

☒ HTTP2

Send requests to targets using HTTP/2. Supported when the request protocol is HTTP/2 or gRPC, but gRPC-specific features are not available.

☐ gRPC

Send requests to targets using gRPC. Supported when the request protocol is gRPC.

Created target group for Elastic load balancer.

Nginx-Load-Balancer

Actions

▼ Details

Load balancer type

Application

Status

Active

VPC

vpc-097825383c424a849

IP address type

IPv4

Scheme

Internet-facing

Hosted zone

Z23TAZ6LKFMNIO

Availability Zones

subnet-0f285dce115e66052 eu-north-1a (eu-north-1-az1)
subnet-0e0c87f9300987cab eu-north-1b (eu-north-1-az2)

Date created

July 20, 2023, 11:20 (UTC+03:00)

Load balancer ARN

arn:aws:elasticloadbalancing:eu-north-1:273737642040:loadbalancer/app/Nginx-Load-Balancer/643800ca1dc94655

DNS name

info
Nginx-Load-Balancer-749436698.eu-north-1.elb.amazonaws.com (A Record)

Listeners and rules

Network mapping

Security

Monitoring

Integrations

Attributes

Tags

Listeners and rules (1) info

Manage rules

Manage listener

Add listener

A listener checks for connection requests on its configured protocol and port. Traffic received by the listener is routed according to the default action and any additional rules.

Filter listeners by property or value

< 1 > ⚙

Protocol:Port

▼

Default action

▼

Rules

▼

ARN

▼

Security policy

▼

Default SSL cert

▼

Tags

▼

HTTP:80

Forward to target group

Nginx-Target-Group: 1 (100%)

Group-level stickiness: Off

1 rule

ARN

Not applicable

Not applicable

0 tags